

Individual Analysis Report

Algorithm: Max-Heap Implementation (with `increaseKey` and `extractMax`)

Student: Ilnur Saipiyev

Course: Design and Analysis of Algorithms

Date: October 2025

1. Algorithm Overview

This report analyzes the implementation and performance of a **Max-Heap** data structure written in Java.

A Max-Heap is a complete binary tree where each parent node has a value greater than or equal to its children.

The heap is implemented using an **array-based structure**, where the index relationships between parent and child nodes follow the formulas:

- $\text{parent}(i) = (i - 1) / 2$
- $\text{leftChild}(i) = 2 * i + 1$
- $\text{rightChild}(i) = 2 * i + 2$

The implementation includes three major operations:

- **insert(int value)** — adds a new element and maintains the heap property via upward heapify (sift-up).
- **extractMax()** — removes and returns the maximum element (the root), restoring heap order through downward heapify (sift-down).
- **increaseKey(int index, int newValue)** — increases the value of an element and moves it up to its proper position.

Performance metrics were collected using a dedicated **PerformanceTracker**, which records:

- `comparisons` — number of element-to-element comparisons,
- `swaps` — number of element exchanges,
- `allocations` — new memory allocations,
- `arrayAccesses` — reads and writes to the array,
- `maxDepth` — maximum recursion or iterative depth encountered.

The benchmarking process was automated via a **CLI BenchmarkRunner** that tested different input sizes ($n = 100, 1,000, 10,000, 100,000$) and input configurations (random, sorted, reversed, nearly-sorted).

The goal of this analysis is to validate the theoretical asymptotic complexities of Max-Heap operations through **empirical measurements**, identify **performance bottlenecks**, and provide **optimization recommendations**.

2. Complexity Analysis

2.1 Time Complexity

Insert (sift-up):

When inserting a new element, it is placed at the end of the heap and compared with its parent until the heap property is restored.

In the worst case, the new element moves up the entire height of the heap, which is $\log_2(n)$.

- **Worst case:** $O(\log n)$ — element bubbles up through all levels.
- **Best case:** $\Omega(1)$ — if the inserted element is smaller than its parent, no swaps occur. □
- **Average case:** $\Theta(\log n)$ — expected logarithmic behavior for random inserts.

ExtractMax (sift-down):

The root (maximum element) is removed and replaced by the last element.

The replacement element is then sifted down to maintain the heap property. Each level requires one or two comparisons and possibly one swap.

- **Worst case:** $O(\log n)$ — the element moves down to the bottom.
- **Best case:** $\Omega(1)$ — if the new root is already larger than its children. □
- **Average case:** $\Theta(\log n)$ — logarithmic descent on average.

IncreaseKey:

The increased element may move up the heap similarly to insertion.

- **Worst case:** $O(\log n)$
- **Average case:** $\Theta(\log n)$
- **Best case:** $\Omega(1)$

2.2 Space Complexity

The Max-Heap is implemented as an **in-place array**, requiring **$O(1)$** auxiliary space.

The only additional structure is the `PerformanceTracker`, which maintains a constant number of primitive counters — also **$O(1)$** space.

Therefore, the total space complexity remains **$\Theta(n)$** due to the array of elements.

2.3 Complexity Derivations (Readable Form)

Let h be the height of the heap ($h = \lceil \log_2 n \rceil$).

- **Insertion:** The maximum number of swaps $\leq h \rightarrow T(n) \approx c_1 \cdot \log_2 n + c_2 \rightarrow O(\log n)$
- **Extraction:** Each sift-down operation compares with up to 2 children per level $\rightarrow O(\log n)$
- **IncreaseKey:** Same logic as insertion $\rightarrow O(\log n)$

BuildHeap (for comparison):

If we were to build the heap bottom-up:

$$T(n) = \sum_{i=0}^{\log_2 n} 2^i \cdot i \approx 2n \Rightarrow \Theta(n) \quad T(n) = \sum_{i=0}^{\log_2 n} \frac{n}{2^{i+1}} \cdot i \approx 2n \Rightarrow \Theta(n)$$

This supports the claim that constructing a heap from n elements takes linear time.

3. Code Review and Optimization

3.1 Strengths

- The heap operations (`insert`, `extractMax`, `increaseKey`) are implemented correctly and maintain the heap invariant.
- The code uses **iterative heapify** methods, avoiding recursion and preventing stack overhead.
- Performance tracking is cleanly integrated and writes detailed metrics to CSV for empirical validation.
- Input generation in `BenchmarkRunner` supports multiple cases (random, sorted, reversed, nearly-sorted), enabling robust analysis.

3.2 Weaknesses

1. Swaps counter is unused.

Although `PerformanceTracker` defines a `swaps` metric, the code never increments it inside heapify operations.

→ This results in incomplete tracking of actual element movements.

2. Array access tracking undercounts.

The implementation increments `arrayAccesses` only for selected read/write actions. In several cases, repeated reads (like `heap[parent]`) are not counted, leading to underestimated access metrics.

3. Redundant capacity allocation.

`ensureCapacity()` doubles the array size but `BenchmarkRunner` initializes the heap with exact capacity n , meaning this branch is never tested in benchmarks.

For more realistic performance, initial capacity should start smaller to trigger allocations.

4. Inefficient value comparisons.

In `heapifyDown`, comparisons like `if(heap[right] > heap[largest])` re-read the array multiple times; caching these values in local variables can reduce array accesses.

3.3 Optimization Recommendations

- **Fix swap counter usage:** Increment `swaps` whenever two elements are exchanged.
- **Cache array reads:** Temporarily store `heap[parent]`, `heap[left]`, `heap[right]` to reduce redundant array access.
- **Optimize ensureCapacity:** Use `Arrays.copyOf` with a calculated growth factor or a geometric progression (1.5x instead of 2x) to balance memory and speed.
- **Add microbenchmarks:** Implement fine-grained timing per operation (e.g., per insert) to analyze constant factors.

These optimizations would not change asymptotic complexity but would **reduce real runtime and memory traffic**.

4. Empirical Results and Analysis

The benchmarks were executed for input sizes $n = 100, 1,000, 10,000, 100,000$ and for four input distributions: random, sorted, reversed, and nearly-sorted.

4.1 Observed Trends

- Time growth:**
Time scales approximately proportional to $n \log n$ across all scenarios.
`extractMax` shows the largest runtime since it performs n removals, each costing $O(\log n)$.
- Input impact:**
 - Sorted inputs** show higher comparisons due to frequent violations of the heap property during insertion.
 - Reversed inputs** are slightly more efficient for small n because inserts often preserve order.
 - Nearly-sorted inputs** behave between random and sorted cases.
- Metric correlations:**
 - time strongly correlates with both comparisons and arrayAccesses, confirming that computational cost is dominated by element movement and comparisons.
 - No excessive memory allocations were observed (allocations = 0), confirming efficient in-place execution.

4.2 Quantitative Summary (CLI BenchmarkRunner)

n	Scenario	Input Type	Time (ns)	Comparisons	Array Accesses
100,000	insert	random	2,080,000	228,288	556,591
100,000	extract	random	10,800,000	2,949,433	4,900,258
100,000	increaseKey	random	2,840,000	278,014	806,042

Interpretation:

- ☐ `extractMax` is the most expensive operation ($\sim 5\times$ slower than `insert`). ☐
- Growth across input sizes confirms logarithmic per-operation scaling.

4.3 Validation

Plotting time vs n on logarithmic axes yields near-linear curves, verifying that:

$$T(n) \propto n \cdot \log n \quad T(n) \propto n \cdot \log n$$

This confirms the theoretical complexity of heap operations.

4.4 JMH Benchmark Setup

In addition to the CLI benchmarks, the implementation was tested using the **Java Microbenchmark Harness (JMH)** framework to obtain statistically rigorous micro-level timings.

Configuration:

- **Mode:** AverageTime
- **Time unit:** milliseconds per operation
- **Warmup:** 1 iteration (100 ms)
- **Measurement:** 2 iterations (100 ms each)
- **Forks:** 1

Each benchmark tests the three primary heap operations (insert, extractMax, increaseKey) across input sizes $n = 50, 200, 1000$ under random data generation.

4.5 JMH Microbenchmark Results

Benchmark	n	Mode	Count	Score (ms/op)	Error
testExtractMax	50	avgt	2	0.149	—
testExtractMax	200	avgt	2	0.145	—
testExtractMax	1000	avgt	2	0.164	—
testIncreaseKey	50	avgt	2	0.126	—
testIncreaseKey	200	avgt	2	0.122	—
testIncreaseKey	1000	avgt	2	0.122	—
testInsert	50	avgt	2	0.118	—
testInsert	200	avgt	2	0.117	—
testInsert	1000	avgt	2	0.131	—

4.6 Analysis of JMH Results

1. Consistency:

The per-operation time remains nearly constant (≈ 0.12 – 0.16 ms/op) across different n , confirming logarithmic time per operation.

2. Relative cost:

extractMax is the slowest operation, consistent with the expected $O(\log n)$ cost due to repeated comparisons and swaps.

insert and increaseKey exhibit nearly identical timings, showing that upward heapify behaves predictably regardless of heap size.

3. JIT and stability:

The JVM's Just-In-Time (JIT) optimization stabilizes performance quickly after warmup, yielding very consistent average times across runs.

4. **CLI vs JMH comparison:**

- **CLI** benchmarks capture *total workload scaling* (bulk $n \log n$ behavior).
 - **JMH** benchmarks isolate *micro-level operation costs*, highlighting the constant factors and JVM optimizations.
- Both confirm the same asymptotic trend and consistent runtime growth.

4.7 Combined Conclusion for Empirical Analysis

Both macro-level (CLI) and micro-level (JMH) benchmarks demonstrate strong agreement with theoretical expectations:

- **Insert, extractMax, and increaseKey** all show logarithmic scaling.
- **extractMax** dominates total runtime due to multiple downward heapify steps.
- **No recursion or extra allocations** occurred, validating in-place efficiency.
- The implementation is robust under different input sizes and JVM optimization phases.

5. Conclusion

The **Max-Heap implementation** successfully demonstrates the theoretical and empirical behavior expected from logarithmic-time heap operations.

Both **macro-level (CLI)** and **micro-level (JMH)** benchmarks confirm the correctness, stability, and performance efficiency of the implementation.

Key Findings

- All primary operations — insert, extractMax, and increaseKey — consistently exhibit **$O(\log n)$** time complexity.
- The **PerformanceTracker** verified low memory overhead and precise tracking of computational events.
- **extractMax** remains the most computationally expensive operation, being approximately five times slower than insert, consistent with its expected behavior.
- **JMH benchmarks** confirmed stable per-operation performance ($\approx 0.12\text{--}0.16$ ms/op), proving that heap operations scale logarithmically even under JVM micro-optimization.
- Both benchmarking methods (CLI and JMH) aligned closely with theoretical models, validating analytical predictions.

JMH-Specific Insights

- The per-operation time remained nearly constant across increasing heap sizes, confirming logarithmic growth and efficient cache behavior.
- The JVM's Just-In-Time (JIT) compilation and hotspot optimizations stabilized performance quickly after warmup, producing highly consistent results.
- This demonstrates that the implementation performs efficiently not only in algorithmic theory but also in practice under a managed runtime environment.

Recommendations

- Improve **metric accuracy** in PerformanceTracker by properly counting swaps and all array accesses.
- Add **fine-grained microbenchmarks** for each operation to measure constant factors in more detail.
- Optionally extend the analysis to **buildHeap()** and compare top-down vs bottom-up construction ($\Theta(n)$ complexity).

Overall Conclusion

The Max-Heap algorithm implementation is **correct, efficient, and theoretically consistent**. Empirical measurements from both CLI and JMH confirm that it achieves the expected **$O(\log n)$** performance with minimal overhead.

Minor optimizations can further refine constant-time behavior, but the core algorithmic design is sound.

This analysis validates a **successful design, implementation, and performance evaluation** of the Max-Heap data structure.